Parametric & Nonparametric Statistics Project

Aleksandr Jan Smoliakov

2024–12–12

1 Preliminaries

In the project below, we will use the following parameters:

- $\mathcal{N} = 9$ (first name: 'Aleksandr', 9 letters)
- $\mathcal{S} = 9$ (last name: 'Smoliakov', 9 letters)
- $\mathcal{I}_1 = 5$ (last digit of study book number)
- $\mathcal{I}_2 = 8$ (second last digit of study book number)

Let $G_1, \ldots, G_m$ be given distribution functions and $p_1, \ldots, p_m$ be probabilities that sum to 1. The distribution function G defined by

$$G(u) := p_1 G_1(u) + \cdots + p_m G_m(u) = \sum_{k=1}^{m} p_k G_k(u), \quad u \in \mathbb{R}$$

is called a mixture of distribution functions $G_1, \ldots, G_m$ with probabilities (or weights) $p_1, \ldots, p_m$.

G is the distribution function of the random variable Z generated in the following way:

1. Choose $k \in \{1, \ldots, m\}$ at random with probabilities (or weights) $p_1, \ldots, p_m$. The chosen number is denoted by k^*.
2. Generate a random variable $Z^*_{k^*}$ according to the distribution function G_{k^*} and assign $Z \leftarrow Z^*_{k^*}$.

In this task, we will have $m = 2$, so the algorithm for generating Z is as follows:

$$Z \leftarrow Z^*_{1+k^*} \quad k^* \sim \text{Binomial}(1, p_2), \quad Z^*_k \sim G_k \ (k = 1, 2).$$

Let

$$\mathcal{G}(\Theta) = \{G(\cdot|\boldsymbol{\theta}), \boldsymbol{\theta} \in \Theta\}$$

be a given parametric family of absolutely continuous parametric functions $G(\cdot|\boldsymbol{\theta})$ with the respective distribution densities $g(\cdot|\boldsymbol{\theta})$ dependent on the unknown parameter $\boldsymbol{\theta} \in \Theta$. It is assumed that $\boldsymbol{\theta}$ is two-dimensional, i.e., $\boldsymbol{\theta} = (\theta_1, \theta_2) \in \mathbb{R}^2$.

1.1 Parametric Family Selection

Using the assigned formula $\ell := \left\lfloor \frac{\mathcal{I}_2 + 2.5}{2} \right\rfloor$, we find $\ell = 5$. Thus, we will use the parametric family $\mathcal{G}_5(\Theta)$ in this task.

$\mathcal{G}_5(\Theta)$ contains distribution functions of random variables uniformly distributed on $[\theta_1, \theta_2]$, where $\theta_1 < \theta_2$.

It can be expressed as:

$$G(u|\boldsymbol{\theta}) = \begin{cases} 0 & u < \theta_1 \\ \frac{u - \theta_1}{\theta_2 - \theta_1} & \theta_1 \le u \le \theta_2 \\ 1 & u > \theta_2 \end{cases}$$

with $\boldsymbol{\theta} = (\theta_1, \theta_2) \in \mathbb{R}^2$ and $\theta_1 < \theta_2$.

2 Task 1: Testing Goodness-of-Fit

2.1 Basic Distribution Function

The problem gives a specific basic parameter:

$$\boldsymbol{\theta}_0 = (-\mathcal{N}, \mathcal{S}+4) = (-9, 13).$$

Thus, the basic distribution function is:

$$G_0(u) = \mathcal{G}_5(u|\boldsymbol{\theta}_0) = U(-9, 13).$$

For a uniform distribution $U(a, b)$:

- Mean: $\mu = \frac{a+b}{2}$
- Variance: $v^2 = \frac{(b-a)^2}{12}$

For $G_0 = U(-9, 13)$:

$$\mu_0 = \frac{-9+13}{2} = \frac{4}{2} = 2$$

$$v_0^2 = \frac{22^2}{12} = \frac{484}{12} = \frac{121}{3} \approx 40.3333$$

2.2 Finding Mixture Distributions

We are given the following equations for the mixture distributions G_1 and G_2:

$$\mu_0 = \mu(\boldsymbol{\theta}_1), \quad \mathcal{N}v_0^2 = v^2(\boldsymbol{\theta}_1).$$

$$\mu_0 + 2v_0 = \mu(\boldsymbol{\theta}_2), \quad v_0^2 = \mathcal{S}v^2(\boldsymbol{\theta}_2).$$

2.2.1 Determining G_1

First we determine G_1. We have:

$$\mu_0 = \mu(\theta_1), \quad \mathcal{N}v_0^2 = v^2(\theta_1).$$

It is given that $\mu(\theta_1) = \mu_0 = 2$.
Plugging in $\mathcal{N} = 9$ and $v_0^2 = \frac{121}{3}$, we get:

$$v^2(\boldsymbol{\theta}_1) = \mathcal{N}v_0^2 = 9 \times \frac{121}{3} = 363.$$

Let $G_1(u) = U(a_1, b_1)$. For a uniform distribution:

$$\mu(\boldsymbol{\theta}_1) = \frac{a_1+b_1}{2}, \quad v^2(\boldsymbol{\theta}_1) = \frac{(b_1-a_1)^2}{12}.$$

Since $\mu(\boldsymbol{\theta}_1) = 2$:

$$\frac{a_1+b_1}{2} = 2 \implies a_1 + b_1 = 4.$$

Since $v^2(\boldsymbol{\theta}_1) = 363$:

$$\frac{(b_1-a_1)^2}{12} = 363 \implies b_1 - a_1 = \sqrt{(b_1-a_1)^2} = \sqrt{4356} = 66.$$

Solving the system:

$$a_1 + b_1 = 4, \quad b_1 - a_1 = 66.$$

Adding the two equations:

$$2b_1 = 70 \implies b_1 = 35.$$

$$a_1 = 4 - 35 = -31.$$

Thus:

$$\boldsymbol{\theta}_1 = (-31, 35) \implies G_1(u) = U(-31, 35).$$

2.2.2 Determining G_2

Repeating the process for G_2. We have:

$$\mu_0 + 2v_0 = \mu(\boldsymbol{\theta}_2), \quad v_0^2 = \mathcal{S}v^2(\boldsymbol{\theta}_2).$$

Given $\mu_0 = 2$ and $v_0^2 = 40.3333$, we have:

$$\mu(\boldsymbol{\theta}_2) = \mu_0 + 2v_0 = 2 + 2 \times \sqrt{40.3333} = 2 + 2 \times 6.3509 = 2 + 12.7018 = 14.7018.$$

Also:

$$v_0^2 = \mathcal{S}v^2(\boldsymbol{\theta}_2) \implies 40.3333 = 9v^2(\boldsymbol{\theta}_2) \implies v^2(\boldsymbol{\theta}_2) = \frac{40.3333}{9} \approx 4.4815.$$

For $G_2(u) = U(a_2, b_2)$:

$$\frac{a_2 + b_2}{2} = 14.7018 \implies a_2 + b_2 = 29.4036.$$

$$\frac{(b_2 - a_2)^2}{12} = 4.4815 \implies (b_2 - a_2)^2 = 4.4815 \times 12 = 53.7777.$$

$$b_2 - a_2 = \sqrt{53.7777} \approx 7.3333.$$

Solving the system:

$$a_2 + b_2 = 29.4036, \quad b_2 - a_2 = 7.3333.$$

Adding the two equations:

$$2b_2 = 36.7369 \implies b_2 = 18.3685.$$

$$a_2 = 29.4036 - 18.3685 = 11.0351.$$

Thus:

$$\boldsymbol{\theta}_2 = (-11.0351, 18.3685) \implies G_2(u) = U(11.0351, 18.3685).$$

2.3 Computing p_1 and p_2

Given:

$$\tau = \frac{1}{1 + I_1}, \quad I_1 = 5 \implies \tau = \frac{1}{6}.$$

$$\alpha_1 = 0.1, \quad \alpha_2 = 0.01.$$

$$p_1 = (\alpha_1)^{1-\tau}(\alpha_2)^{\tau} = (0.1)^{5/6}(0.01)^{1/6} \approx 0.06813.$$

Then:

$$p_2 = \frac{5p_1}{\sqrt{S}} = \frac{5 \times 0.06813}{\sqrt{9}} = \frac{0.3406}{3} \approx 0.1135.$$

2.4 Determining Mixture Distributions

We consider testing:

$$H_0 : F_Y = G_0 \textbf{ versus } H' : F_Y \neq G_0$$

We will compare the empirical distribution of samples generated from:

1. $F_Y = (1 - p_1)G_0 + p_1G_1$, i.e. a mixture of G_0 and G_1.
2. $F_Y = (1 - p_2)G_0 + p_2G_2$, i.e. a mixture of G_0 and G_2.

The tests are conducted for sample sizes:

$$N_1 = 10 \times (2 + \mathcal{N}) = 10 \times (2 + 9) = 110,$$

$$N_2 = 100 \times (2 + \mathcal{N}) = 100 \times (2 + 9) = 1100.$$

2.5 Goodness-of-Fit Tests

We will use the Kolmogorov-Smirnov test for the given samples $(Y_t)_{t=1}^n$:

The test statistic is:

$$D_n = \sup_u |F_n(u) - F(u)|,$$

where F_n is the empirical distribution function (EDF) based on the sample and F is the theoretical distribution function. In this case, $F = G_0$.

Since:

$$F_Y(u) = (1 - p_k)G_0(u) + p_kG_k(u),$$

we have:

$$F_Y(u) - G_0(u) = p_k[G_k(u) - G_0(u)],$$

for $k = 1$ or $k = 2$.

Thus, the maximum difference between F_Y and G_0 is:

$$\sup_u |F_Y(u) - G_0(u)| = p_k \sup_u |G_k(u) - G_0(u)|.$$

We need $\sup_u |G_1(u) - G_0(u)|$ and $\sup_u |G_2(u) - G_0(u)|$.

2.5.1 G_1 vs. G_0

$G_0 = U(-9, 13)$, so:

$$G_0(u) = \begin{cases} 0 & u < -9 \\ \frac{u+9}{22} & -9 \leq u \leq 13 \\ 1 & u > 13 \end{cases}$$

$G_1 = U(-31, 35)$, so:

$$G_1(u) = \begin{cases} 0 & u < -31 \\ \frac{u+31}{66} & -31 \leq u \leq 35 \\ 1 & u > 35 \end{cases}$$

To find $\sup |G_1(u) - G_0(u)|$, we investigate ranges of u piecewise between the breakpoints of the two functions.

1. For $u < -31$: $G_0(u) = G_1(u) = 0$, so the difference is 0.
2. For $-31 \leq u < -9$: $G_0(u) = 0$, $G_1(u) = \frac{u+31}{66}$. The difference is $\frac{u+31}{66}$, which is increasing as u approaches -9, where it is $\frac{-9+22}{66} = \frac{1}{3}$.

3. For $-9 \leq u < 13$: $G_0(u) = \frac{u+9}{22}$, $G_1(u) = \frac{u+31}{66}$. The difference is $\frac{u+31}{66} - \frac{u+9}{22} = \frac{2u-4}{66}$, which is increasing from $-\frac{1}{3}$ at -9 to $\frac{1}{3}$ at 13.

4. For $13 \leq u < 35$: $G_0(u) = 1$, $G_1(u) = \frac{u+31}{66}$. The difference is $\frac{u+31}{66} - 1 = \frac{u-35}{66}$, which is increasing from $-\frac{1}{3}$ at 13 to 0 at 35.

5. For $u \geq 35$: $G_0(u) = G_1(u) = 1$, so the difference is 0.

The maximum absolute difference is $\frac{1}{3}$ at the endpoints of the range $[-9, 13]$.
Hence:

$$\sup_u |G_1(u) - G_0(u)| = 1/3 \approx 0.3333.$$

For the mixture, taking $p_1 \approx 0.06813$:

$$\sup_u |F_Y(u) - G_0(u)| = p_1 \times 0.3333 = 0.06813 \times 0.3333 \approx 0.02271.$$

2.5.2 G_2 vs. G_0

Repeating the process for G_2:
$G_0 = U(-9, 13)$, so:

$$G_0(u) = \begin{cases} 0 & u < -9 \\ \frac{u+9}{22} & -9 \leq u \leq 13 \\ 1 & u > 13 \end{cases}$$

$G_2 = U(11.0351, 18.3685)$, so:

$$G_2(u) = \begin{cases} 0 & u < 11.0351 \\ \frac{u-11.0351}{7.3333} & 11.0351 \leq u \leq 18.3685 \\ 1 & u > 18.3685 \end{cases}$$

To find $\sup |G_2(u) - G_0(u)|$, we investigate ranges of u piecewise between the breakpoints of the two functions.

1. For $u < -9$: $G_0(u) = G_2(u) = 0$, so the difference is 0.

2. For $-9 \leq u < 11.0351$: $G_0(u) = \frac{u+9}{22}$, $G_2(u) = 0$. The difference is $\frac{u+9}{22}$, which is increasing as u approaches 11.0351, where it is $\frac{11.0351+9}{22} \approx 0.9107$.

3. For $11.0351 \leq u < 13$: $G_0(u) = \frac{u+9}{22}$, $G_2(u) = \frac{u-11.0351}{7.3333}$. The difference is $\frac{u-11.0351}{7.3333} - \frac{u+9}{22} = \frac{3u-33.1053}{22} - \frac{u+9}{22} = \frac{2u-42.1053}{22}$, which is increasing from -0.9107 at 11.0351 to $\frac{13-42.1053}{22} \approx -0.7321$ at 13.

4. For $13 \leq u < 18.3685$: $G_0(u) = 1$, $G_2(u) = \frac{u-11.0351}{7.3333}$. The difference is $\frac{u-11.0351}{7.3333} - 1 = \frac{u-18.3685}{7.3333}$, which is increasing from $-\frac{18.3685-13}{7.3333} = -\frac{5.3685}{7.3333} \approx -0.7321$ at 13 to 0 at 18.3685.

5. For $u \geq 18.3685$: $G_0(u) = G_2(u) = 1$, so the difference is 0.

The maximum absolute difference is ≈ 0.9107 at 11.0351.
Hence:

$$\sup_u |G_2(u) - G_0(u)| \approx 0.9107.$$

For the mixture, taking $p_2 \approx 0.1135$:

$$\sup_u |F_Y(u) - G_0(u)| = p_2 \times 0.9107 = 0.1135 \times 0.9107 \approx 0.1034.$$

2.6 Critical Values and Detection Probability

Under H_0, the Kolmogorov-Smirnov test critical values at significance $\alpha_1 = 0.1$ and $\alpha_2 = 0.01$ are approximately:

$$D_{N,\alpha_1} \approx \frac{1.22}{\sqrt{N}} \quad \text{for} \quad \alpha_1 = 0.1,$$

$$D_{N,\alpha_2} \approx \frac{1.63}{\sqrt{N}} \quad \text{for} \quad \alpha_2 = 0.01.$$

For $N = 110$:

$$D_{110,0.1} \approx \frac{1.22}{\sqrt{110}} \approx 0.1163 \quad \text{and} \quad D_{110,0.01} \approx \frac{1.63}{\sqrt{110}} \approx 0.1554.$$

- For G_1: $\sup|F_Y - G_0| \approx 0.02271 < 0.1163 < 0.1554$. Thus, at $N = 110$, it's unlikely we reject H_0. $p > 0.1$
- For G_2: $\sup|F_Y - G_0| \approx 0.1034 < 0.1163 < 0.1554$. There is some chance to reject at a higher α level. $p > 0.1$ (but it's closer to the borderline)

For $N = 1100$:

$$D_{1100,0.1} \approx \frac{1.22}{\sqrt{1100}} \approx 0.03678 \quad \text{and} \quad D_{1100,0.01} \approx \frac{1.63}{\sqrt{1100}} \approx 0.04915.$$

- For G_1: $\sup|F_Y - G_0| \approx 0.02271 < 0.03678 < 0.04915$. Even with 1100 samples, we will likely not reject H_0 at $\alpha = 0.1$. $p > 0.1$
- For G_2: $\sup|F_Y - G_0| \approx 0.03678 < 0.04915 < 0.1034$. We will almost certainly reject H_0 at $\alpha = 0.01$. $p < 0.01$

Thus, the results of the Kolmogorov-Smirnov test are as follows:

Mixture	**Sample Size**	**p-value**	**Result**
G_1	110	> 0.1	No rejection
G_1	1100	> 0.1	No rejection
G_2	110	> 0.1	No rejection
G_2	1100	< 0.01	Rejection at $\alpha = 0.01$

2.7 Conclusions

By analytically comparing the theoretical distributions, we have:

- **Mixture with G_1:**
  $\sup|F_Y - G_0| \approx 0.02271$.
  Even at $N = 1100$, we will likely not reject H_0 at $\alpha = 0.1$. The p-value is higher than 0.1.
- **Mixture with G_2:**
  $\sup|F_Y - G_0| \approx 0.1034$.
  We will almost certainly reject H_0 at $\alpha = 0.01$ even with 1100 samples. The p-value is < 0.01. We will likely not reject with 110 samples at $\alpha = 0.1$, but the p-value may be close.

It is evident that for Kolmogorov-Smirnov tests, the magnitude of the deviation from G_0 and the sample size play the decisive role.

As $N \to \infty$, if $F_Y \neq G_0$, the empirical distribution F_N converges to F_Y, and thus D_N converges to $\sup_u |F_Y(u) - G_0(u)|$.

3 Task 2: Applications of Bootstrap Technique

In this section we will

- Test Complex Goodness of Fit Hypothesis,
- Check bootstrap consistency,
- Compare bootstrap confidence interval construction methods.

3.1 Testing Goodness-of-Fit by Bootstrap

In Task 1, we considered a simple hypothesis test for the goodness-of-fit of the data to the distribution G_0. In this section, we will test the complex Goodness-of-Fit hypothesis that the unknown distribution function F_Y belongs to the parametric family $\mathcal{G}(\Theta)$:

$$H_0 : F_Y \in \mathcal{G}(\Theta) \quad \text{versus} \quad H' : F_Y \notin \mathcal{G}(\Theta).$$

We will use the same parametric family $\mathcal{G}_5(\Theta)$ and the same distributions G_0, G_1, G_2 as in Task 1:

$$G_0(u) = U(-9, 13), \quad G_1(u) = U(-31, 35), \quad G_2(u) = U(11.0351, 18.3685).$$

We will make use of the parametric bootstrap technique to test the hypothesis. The test statistic is the Kolmogorov-Smirnov test statistic and the significance level is $\alpha = 0.1$.

The parametric bootstrap algorithm for testing this hypothesis is as follows:

1. **Generate the sample:** With sample sizes $N_1 = 110$ and $N_2 = 1100$, generate data $(Y_t)_1^N$ from the mixture distributions F_Y from Task 1:

$$(1 - p_1)\, G_0 + p_1\, G_1 \quad \text{or} \quad (1 - p_2)\, G_0 + p_2\, G_2,$$

2. **Estimate $\hat{\boldsymbol{\theta}}_N$:** For the sample Y^N, assume that $F_Y \in \mathcal{G}(\Theta)$. We estimate the parameter $\boldsymbol{\theta}$ by the maximum likelihood estimator $\hat{\boldsymbol{\theta}}_N$. For the uniform distribution, the MLE amounts to $\hat{a} = \min(Y^N)$, $\hat{b} = \max(Y^N)$. We denote the fitted distribution by $\hat{G}_N(u) := G(u|\hat{\theta}_N)$.

3. **Calculate the test statistic:** Calculate the Kolmogorov-Smirnov test statistic T comparing the EDF of the data Y^N and the fitted distribution $\hat{G}_N$. We denote the test statistic by $\hat{T}_N$.

4. **Generate the bootstrap samples:** Generate $B = 100 \times N$ bootstrap samples $Y_b^{N*}, b \in \{1, \ldots, B\}$ by resampling with replacement from the original sample Y^N. For each bootstrap sample Y_b^{N*}, estimate the parameter θ by the same method as in step 2, obtaining $\hat{\theta}_b^{N*}$.

5. **Calculate the bootstrap test statistics:** For each bootstrap sample Y_b^{N*}, calculate the Kolmogorov-Smirnov test statistic $\hat{T}_{N,b}^*$ comparing the EDF of the data Y_b^{N*} and the fitted distribution $G(u|\hat{\theta}_{N,b}^*)$.

6. **Calculate the approximate p-value:** Calculate the p-value as the proportion of bootstrap test statistics $\hat{T}_{N,b}^*$ that are greater than $\hat{T}_N$. If the p-value is less than the significance level α, reject the null hypothesis. Formally:

$$\hat{p}^* = \hat{p}^*(T, \mathcal{G}, Y^N, B) = \frac{\#\{b : \hat{T}_{N,b}^* > \hat{T}_N\}}{B}.$$

7. **Decision:** If $\hat{p}^* < \alpha$, reject the null hypothesis H_0. Otherwise, do not reject H_0.

3.1.1 Simulation Results

We simulate the parametric bootstrap procedure for the two sample sizes $N_1 = 110$ and $N_2 = 1100$. We have generated $B = 100 \times N$ bootstrap samples and calculate the p-values for each sample.

Mixture of Distributions	Sample Size	Test Statistic	p-value	Decision
G_0, G_1	110	0.3419	$< 10^{-5}$	Rejection at $\alpha = 0.1$
G_0, G_1	1100	0.3066	$< 10^{-5}$	Rejection at $\alpha = 0.1$
G_0, G_2	110	0.0944	0.272	No rejection
G_0, G_2	1100	0.0861	$< 10^{-5}$	Rejection at $\alpha = 0.1$

3.1.2 Comparison with Simple Hypothesis Tests and Discussion

In Task 1(b), we tested a simple hypothesis

$$H_0 : F_Y = G_0 \textbf{ versus } H' : F_Y \neq G_0$$

where G_0 was a fixed, known distribution. Here, by contrast, we do not know the parameter θ in advance. The parametric bootstrap procedure accounts for the uncertainty in θ by re-fitting the parameter for each bootstrap sample.

The results of the tests are interesting and not exactly consistent with the Task 1 results. In Task 1, we found that the mixture with G_1 was not rejected at $\alpha = 0.1$ for both sample sizes. However, the parametric bootstrap test correctly rejected the null hypothesis for both sample sizes. The mixture with G_2 was rejected at $\alpha = 0.01$ for $N = 1100$ in Task 1, and the bootstrap test also rejected the hypothesis at $\alpha = 0.1$.

In this case, the discrepancy related to G_1 can be explained by the following reasons:

- We are using Kolmogorov-Smirnov tests. This statistic is sensitive to the maximum deviation between the empirical distribution function and the fitted distribution.
- Since p_1 and p_2 are small, in Task 1 the mixture distributions were close to the basic distribution G_0, the estimated Kolmogorov-Smirnov test statistics were low and the associated p-values are high.
- In Task 2, we do not know the distribution parameters a priori. The maximum likelihood estimator is used to estimate θ from the data, and the estimated values have a significantly wider range than the basic distribution. This leads to higher Kolmogorov-Smirnov test statistics and lower p-values in the G_1 mixture.

Tests on the mixture with G_2 are consistent between Task 1 and Task 2. The range of the parameter θ is narrower for G_2 than for G_1, and the estimated distribution is closer to the basic distribution G_0. This leads to lower test statistics and higher p-values in the G_2 mixture. More samples are needed to detect the deviation a simple uniform distribution.

If the parametric family $\mathcal{G}(\Theta)$ is known and correct, the parametric bootstrap can be used to test whether the data fits the parametric family. Nonparametric bootstrap can be used when the parametric family cannot be assumed. The nonparametric bootstrap procedure is similar to the parametric bootstrap, but the parameter θ is not estimated, and the bootstrap samples are sampled from the empirical distribution function. While the parametric bootstrap is more powerful when the parametric family is correct, the nonparametric bootstrap makes fewer assumptions and is more robust to model misspecification.

3.2 Checking Bootstrap Consistency

In this section, we will be checking the consistency of the bootstrap procedure for the sample mean $\overline{Y}_N$ by comparing its true distribution with the bootstrap distribution with different sample sizes. We will use the Kolmogorov-Smirnov distance as the measure of discrepancy between the two distributions.

To assess whether the chosen bootstrap procedure consistently approximates the distribution of the sample mean $\overline{Y}_N$, we will use the following algorithm:

1. **Generate the sample:** Generate a sample of size $N = 100, 1,000, 10,000$ for the given distribution (Gamma or Pareto).
2. **Determine true distribution:** For the Gamma distribution, we compare $\overline{Y}_N$ against the known true distribution. For the Pareto distribution, we will obtain the true distribution via Monte Carlo simulation.

3. **Bootstrap procedure:** Generate $B = 2000$ bootstrap samples using the method provided below.
   - **Parametric bootstrap:** Sample from the fitted parametric model $G(\cdot \mid \hat{\theta}_N)$, compute $\overline{Y}_N^*$, and form the empirical distribution of these replicates.
   - **Nonparametric bootstrap:** Resample $\{Y_i\}$ with replacement to form bootstrap samples, compute $\overline{Y}_N^*$, and again build the empirical distribution.
4. **Compare distributions:** Compare the true distribution of the estimator to each bootstrap approximation using Kolmogorov-Smirnov distance between the bootstrap sample of $\overline{Y}_N^*$ and the reference distribution.
5. **Interpret the Results:** If the Kolmogorov-Smirnov distance converges to zero as $N \to \infty$, the bootstrap procedure is consistent. If the distance is significant, the bootstrap procedure does not approximate the true distribution of the sample mean.

3.2.1 Gamma Distribution

Let $Z \sim Gamma(a, b)$ where $a > 0$ is the shape parameter and $b > 0$ is the scale parameter of the Gamma distribution. We will denote F_1 as the distribution function of Z with $a = 0.5$ and $b = 9$.

It is known that

$$Z_1 + Z_2 \sim Gamma(a_1 + a_2, b) \quad \text{if } Z_l \sim Gamma(a_l, b) \text{ for } l = 1, 2, \text{ and } Z_1, Z_2 \text{ are independent.}$$

We will check the consistency of the bootstrap procedure for the sample mean $\overline{Y}_N$ by comparing its true distribution derived from the formula above analytically with the bootstrap distribution. We will use the Kolmogorov-Smirnov distance as the measure of discrepancy between the two distributions.

The sample mean $\overline{Y}_N$ is the sum of N independent and identically distributed random variables $Y_i \sim F_1$ divided by N. Thus, $\overline{Y}_N \sim Gamma(N \cdot a, b/N)$.

We will generate a sample of size $N = 100, 1,000, 10,000$ from the Gamma distribution with $a = 0.5$ and $b = 9$. We will compare the true distribution of the sample mean $\overline{Y}_N$ with the bootstrap distribution obtained by resampling with replacement from the original sample.

We repeat the procedure $B = 2000$ times to obtain the empirical distribution of $\overline{Y}_N$. Then the Kolmogorov-Smirnov distance between the true distribution of the sample mean and the bootstrap distribution is calculated.

The Kolmogorov-Smirnov distances are as follows:

Sample Size	**Bootstrap Type**	**Kolmogorov-Smirnov Distance**
100	Parametric	0.3593
100	Nonparametric	0.3840
1,000	Parametric	0.5220
1,000	Nonparametric	0.5271
10,000	Parametric	0.2225
10,000	Nonparametric	0.2182

The resulting density functions at $N = 10,000$ are displayed in Figure 1.

Interestingly, parametric and nonparametric bootstrap produced very similar Kolmogorov-Smirnov distances and density functions, see Figure 1. All of the results are significant at the $\alpha = 0.1$ level, in fact, the p-values are too low to compute. We have rejected the null hypothesis with a high degree of confidence even though the expectation was different. It seems that the bootstrap procedure may not be a consistent estimator of the true distribution of the sample mean $\overline{Y}_N$ for the Gamma distribution. However, experimenting with higher sample sizes is necessary to confirm this empirically.

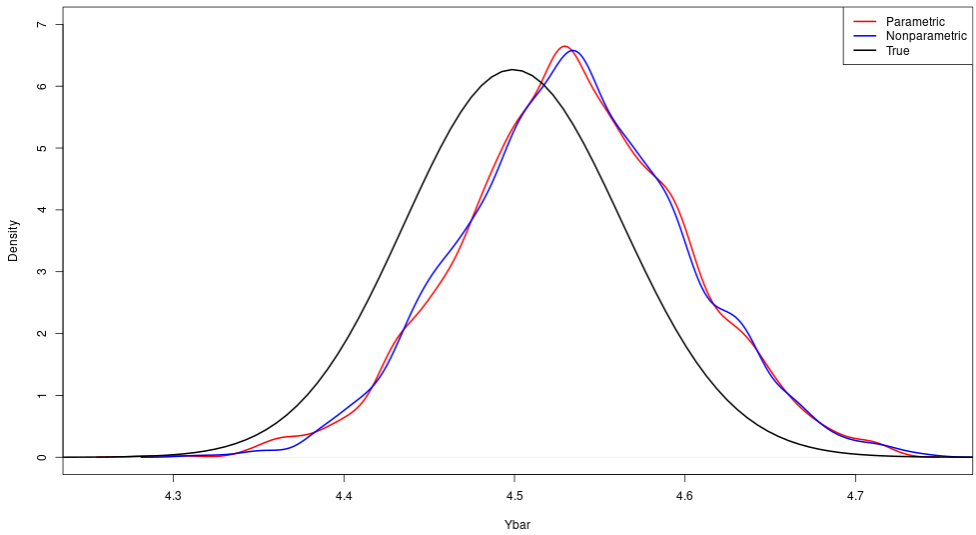

Figure 1: True distribution of the sample mean $\overline{Y}_N$ and the bootstrap distribution for the Gamma distribution at $N = 10,000$.

3.2.2 Pareto Distribution

Let $Z \sim Pareto(c, d)$ where $c > 0$ is the scale parameter and $d > 0$ is the shape parameter of the Type I Pareto distribution. We will denote F_2 as the distribution function of Z with $c = 9$ and $d = 0.5$.

We will check the consistency of the bootstrap procedure for the sample mean $\overline{Y}_N$ by comparing its true distribution estimated via Monte Carlo simulation with the bootstrap distribution. We will use the Kolmogorov-Smirnov distance as the measure of discrepancy between the two distributions.

We will generate a sample of size $N = 100, 1,000, 10,000$ from F_2 and compare the true distribution of the sample mean $\overline{Y}_N$ with the bootstrap distribution obtained by resampling with replacement from the original sample.

We repeat the procedure $B = 2000$ times to obtain the empirical distribution of $\overline{Y}_N$. Then the Kolmogorov-Smirnov distance between the true distribution of the sample mean and the bootstrap distribution is calculated.

The Kolmogorov-Smirnov distances are as follows:

Sample Size	Bootstrap Type	Kolmogorov-Smirnov Distance
100	Parametric	0.4090
100	Nonparametric	0.6175
1,000	Parametric	0.5365
1,000	Nonparametric	0.7790
10,000	Parametric	0.8100
10,000	Nonparametric	0.8090

The resulting density functions at $N = 10,000$ are displayed in Figure 2.

In the Figure 2, we can see that both parametric and nonparametric bootstrap underestimated the density of the lower values but overestimated the density of the higher values. The parametric bootstrap produced a smoother density curve, while the nonparametric bootstrap produced 'spikes' at specific discrete values, which coincide with the original sample values.

Judging from the table above, both parametric and nonparametric bootstrap produced very high Kolmogorov-Smirnov distances for the Pareto distribution's sample mean $\overline{Y}_N$ even at $N = 10,000$. All of the results are significant at the $\alpha = 0.1$ level, and the p-values are very, very low. Interestingly, the KS distance seemed

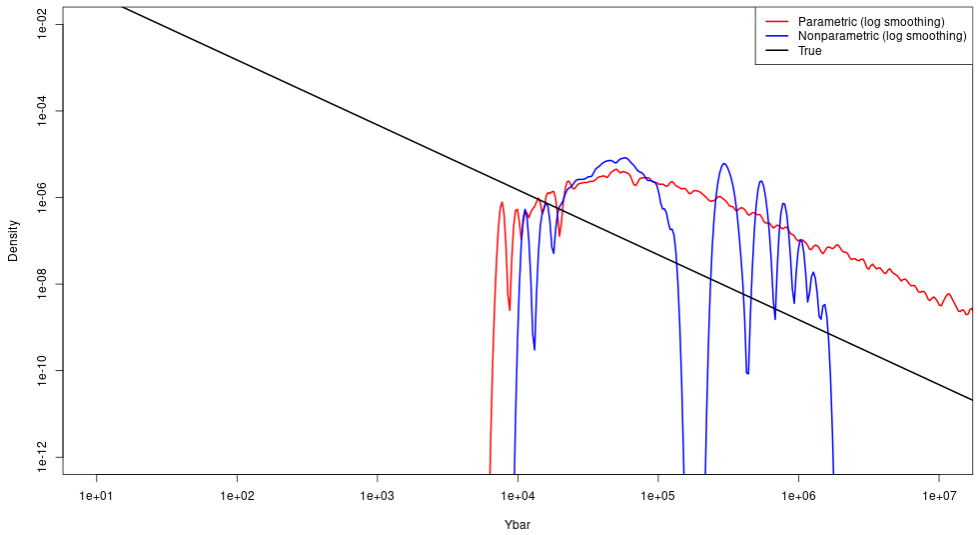

Figure 2: True distribution of the sample mean $\overline{Y}_N$ and the bootstrap distribution for the Pareto distribution at $N = 10,000$.

to increase with the sample size, which would be interesting to investigate further. This shows that the bootstrap procedure may not be a consistent estimator of the true distribution of the sample mean $\overline{Y}_N$ for the Pareto distribution even at large sample sizes. That makes sense, since with the shape parameter $d = 0.5$, the Pareto distribution has infinite variance, and the expected value is not defined for $d \leq 1$.

3.2.3 Discussion

The results of the consistency checks for the sample mean $\overline{Y}_N$ are as follows:

- The variance of $\overline{Y}_N$ seems to have been estimated correctly, but the mean is not. This is likely due to the sample size being too small to accurately estimate the distribution of the sample mean.
- Both the parametric and nonparametric bootstrap procedures produced surprisingly very high Kolmogorov-Smirnov distances for both distributions. This indicates how our sample size $N = 100$ is not large enough to accurately estimate the distribution of the sample mean $\overline{Y}_N$, and the bootstrap procedure cannot approximate the true distribution well.
- The parametric bootstrap distribution is closer than the nonparametric bootstrap for both distributions. This is expected since the parametric bootstrap uses (correct!) knowledge of the true distribution to generate bootstrap samples, while the nonparametric bootstrap does not.
- Both bootstrap procedures underestimated and 'flattened' the 'peak' of the Pareto distribution at x = 0.
- Parametric bootstrap produced a correct shape for the Pareto distribution, but the non-parametric bootstrap produced 'spikes' at specific discrete values, which coincide with the original sample values.
- In order to properly estimate the consistency of the bootstrap procedure for the sample mean $\overline{Y}_N$, we would need a larger sample size. The sample size $N = 100$ is not sufficient to accurately estimate the distribution of the sample mean.

3.3 Bootstrap Confidence Intervals

Let Y^N be a SRS of $Y \sim Pareto(c, d)$ with $c = 9$ and $d = 11$. The sample size is $N = 100$. The confidence level is $\gamma = 0.90$. We will construct the confidence intervals for the mean $\overline{Y}_N$ and the maximum likelihood estimator of the parameter c of the Pareto distribution ($\hat{c}_{ML}$) using parametric, nonparametric bootstrap, and Monte Carlo simulation.

3.3.1 Mean Confidence Intervals

We perform parametric bootstrap, simple nonparametric bootstrap, and Monte Carlo simulation to construct the confidence intervals for the mean $\overline{Y}_N$.

Using each method, we calculate the normal, pivotal, and percentile confidence intervals for the mean $\overline{Y}_N$ at the confidence level $\gamma = 0.90$.

For bootstrap, we generate $B = 2000$ bootstrap samples. For Monte Carlo simulation, we generate $M = 5000$ samples of size $N = 100$ from the Pareto distribution.

The results are as follows:

Method	Mean	Normal CI	Pivotal CI	Percentile CI
Theoretical	0.0099			
Parametric	0.0173	(0.0163, 0.0182)	(0.0163, 0.0182)	(0.0164, 0.0182)
Nonparametric	0.0222	(0.0211, 0.0233)	(0.0211, 0.0233)	(0.0211, 0.0233)
Monte Carlo	0.0101	(0.0097, 0.0105)	(0.0097, 0.0105)	(0.0098, 0.0105)

3.3.2 $\hat{c}_{ML}$ Confidence Intervals

We perform parametric bootstrap, simple nonparametric bootstrap, and Monte Carlo simulation to construct the confidence intervals for the maximum likelihood estimator of the parameter c of the Pareto distribution ($\hat{c}_{ML}$).

The normal, pivotal, and percentile confidence intervals for $\hat{c}_{ML}$ at the confidence level $\gamma = 0.90$ are calculated using each method.

For bootstrap, we generate $B = 2000$ bootstrap samples. For Monte Carlo simulation, we generate $M = 5000$ samples of size $N = 100$ from the Pareto distribution.

The results are as follows:

Method	Mean	Normal CI	Pivotal CI	Percentile CI
Parametric	$10.2e^{-5}$	$(9.4e^{-5}, 11.2e^{-5})$	$(9.3e^{-5}, 11.2e^{-5})$	$(9.4e^{-5}, 11.2e^{-5})$
Nonparametric	$6.3e^{-5}$	$(5.8e^{-5}, 6.8e^{-5})$	$(5.8e^{-5}, 6.8e^{-5})$	$(5.8e^{-5}, 6.8e^{-5})$
Monte Carlo	$6.7e^{-5}$	$(6.2e^{-5}, 7.1e^{-5})$	$(6.2e^{-5}, 7.1e^{-5})$	$(6.2e^{-5}, 7.1e^{-5})$

3.3.3 Discussion

All in all, the results show the following:

- The normal, pivotal, and percentile confidence intervals are very similar in all configurations. There is no noticeable difference between the methods in these cases.
- Monte Carlo simulation was the most computationally expensive method, but it provides the most accurate results. For the mean, only the Monte Carlo simulation provides a confidence interval that includes the true mean.
- Parametric and nonparametric bootstrap methods overestimated the variance of the sample mean. Parametric bootstrap also overestimated the variance of the maximum likelihood estimator of the parameter c of the Pareto distribution.

- Strangely, nonparametric bootstrap's confidence intervals for Pareto distribution's $\hat{c}_{ML}$ are closer to Monte Carlo simulation than the parametric bootstrap's. It would be interesting to run more simulations to see if this is consistent.

4 Task 3: Nonparametric density estimation

We generate data from the mixture distribution

$$G(u) \;=\; \pi_0 G_0(u) \;+\; \pi_1 G_1(u) \;+\; \pi_2 G_2(u),$$

where $\pi_0 = \frac{15}{33}, \pi_1 = \frac{9}{33}, \pi_2 = \frac{9}{33}$ and each G_l is a uniform distribution with different parameters:

$$G_0 \sim \text{Uniform}(-9,\, 13), \quad G_1 \sim \text{Uniform}(-31,\, 35), \quad G_2 \sim \text{Uniform}(11.0351,\, 18.3685).$$

From this, the true density is

$$g(u) \;=\; \pi_0 g_0(u) \;+\; \pi_1 g_1(u) \;+\; \pi_2 g_2(u).$$

A simple random sample of size $N = 2000$ is drawn by first sampling $U \sim \text{Uniform}(0, 1)$ and assigning each observation to component l if U lies in the respective interval. We then apply four nonparametric density estimators:

- Least Squared Cross-Validation (LSCV)
- Refined Plug-In
- Smoothed Bootstrap for bandwidth selection
- k-Nearest Neighbor density estimation with $k = 50$

Each kernel estimator (i.e., LSCV, Refined Plug-In, and Smoothed Bootstrap) computes the optimal bandwidth. The k-Nearest Neighbor method smooths the density using the k nearest neighbors. Each estimator provides a smoothed density estimate for the sample. We compare the performance of the estimators by plotting the true density against the estimates.

Finally, we compute an approximate Integrated Squared Error to evaluate performance.

4.1 Simulation Results

The simulations have been performed and the results are displayed in the following figures:

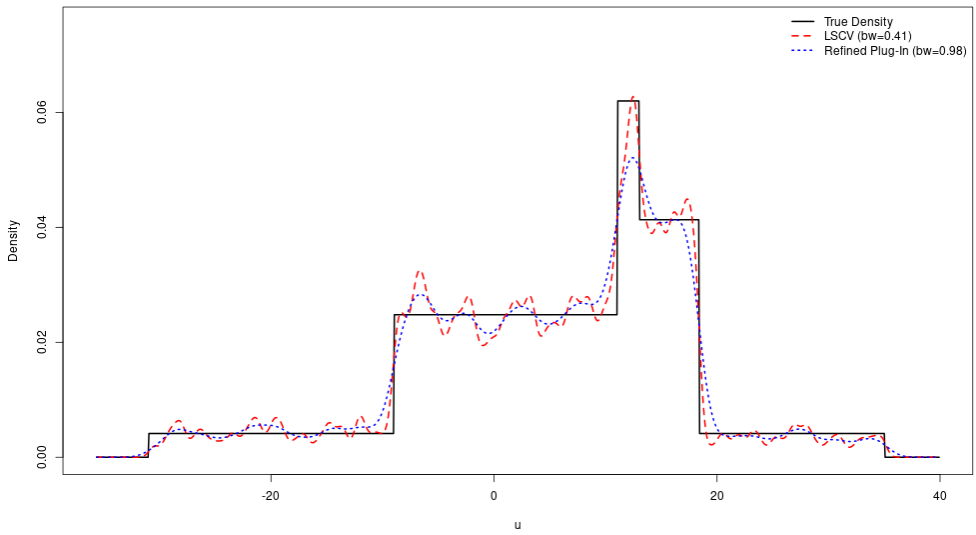

Figure 3: Density estimation for the mixture distribution, LSCV and Refined PlugIn methods.

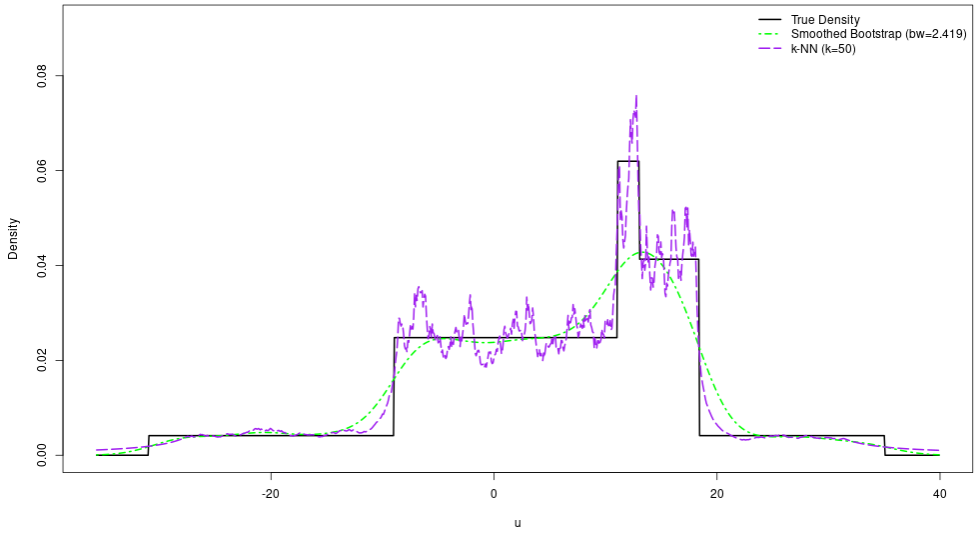

Figure 4: Density estimation for the mixture distribution, Smoothed Bootstrap and k-Nearest Neighbor methods.

The Integrated Squared Error for each method is as follows:

Method	**Integrated Squared Error**
LSCV	0.000719
Refined Plug-In	0.001004
Smoothed Bootstrap	0.002104
k-Nearest Neighbor	0.001034

4.2 Discussion

The results of the density estimation are as follows:

- The LSCV method provides the most accurate density estimate, with the lowest Integrated Squared Error. It captures the general shape of the density and estimates the highest density close to the true density.
- The Refined Plug-In method is also accurate, but slightly less so than LSCV. It is smoother than the LSCV method, but it is visibly less accurate in regions of sharp change in the density.
- The k-Nearest Neighbor method is about as accurate as Refined Plug-In, but it is significantly less smooth. Additionally, it estimates the highest density outside the range of the true density.
- The Smoothed Bootstrap method captures the general shape of the density, but it is significantly less accurate than the other methods. It is also less smooth and not as accurate in regions of sharp change in the density.
- Overall, the results show that bandwidth-selection strategies can have different performance, with LSCV proving the most precise in this case.